

Coursework 2 - Report

Iva Jorgusheska, UID: 11114620

March 28, 2025

1 Harris Corner Detector Function

Corner detection is a fundamental step in image processing and computer vision, especially for identifying keypoints that remain stable across different images of the same scene. Corners are highly distinctive because they exhibit significant changes in intensity along two perpendicular directions, making them robust against translation and rotation.

In this implementation, I developed a Harris Corner Detector in Python, following the standard approach outlined in the lecture slides. The algorithm computes the Harris matrix M at each pixel, applies a Gaussian weighting function, and extracts strong interest points based on the Harris response function R .

1.1 Computation of Image Gradients Using Sobel Operator and Reflection Padding

To compute the partial derivatives of the image with respect to x and y , I used the 3×3 Sobel operator, which approximates the gradient:

$$I_x = \frac{\partial I}{\partial x}, \quad I_y = \frac{\partial I}{\partial y} \quad (1)$$

These gradients are critical because, in a corner, intensity changes in both directions.

The Sobel operator was applied using OpenCV's `cv2.Sobel()` function. Reflection padding (`cv2.BORDER_REFLECT`) was used to prevent artificial border artifacts. The Sobel operator detects edges by computing differences between neighboring pixel intensities, making it ideal for detecting abrupt changes that indicate corners.

```
1 import cv2
2
3 # Compute image gradients using Sobel operator
4 Ix = cv2.Sobel(gray, cv2.CV_64F, 1, 0, ksize=3, borderType=cv2.BORDER_REFLECT)
5 Iy = cv2.Sobel(gray, cv2.CV_64F, 0, 1, ksize=3, borderType=cv2.BORDER_REFLECT)
```

Once the gradients were computed, their squared and cross-product terms were derived to construct the Harris matrix. These terms capture the structure of the local image patches, allowing us to analyze intensity variations in both directions.

```
1 # Compute squared and cross-product gradient terms
2 Ixx = Ix ** 2
3 Iyy = Iy ** 2
4 Ixy = Ix * Iy
```

1.2 Computation of Harris Response R Using a Gaussian Window

To enhance corner localization and suppress noise, the second-moment matrix terms (S_{xx}, S_{yy}, S_{xy}) were computed using a 5×5 Gaussian window with $\sigma = 0.5$. The Gaussian filter smooths the gradient products, reducing noise and making the detector more stable.

```
1 # Apply Gaussian blur to compute second-moment matrix terms
2 sigma = 0.5
3 gaussian_filter = (5, 5)
4 Sxx = cv2.GaussianBlur(Ixx, gaussian_filter, sigma)
5 Syy = cv2.GaussianBlur(Iyy, gaussian_filter, sigma)
6 Sxy = cv2.GaussianBlur(Ixy, gaussian_filter, sigma)
```

Gaussian smoothing ensures that corner detection is robust to small image variations and noise, making it more reliable.

The Harris response R is computed as:

$$R = \det(M) - k(\text{trace}(M))^2 \quad (2)$$

where $k = 0.05$ is an empirically chosen constant. R determines the likelihood of a pixel being a corner.

```

1 # Compute determinant and trace of the matrix M
2 detM = (Sxx * Syy) - (Sxy ** 2)
3 traceM = Sxx + Syy
4
5 # Compute Harris response function
6 R = detM - k * (traceM ** 2)

```

To facilitate thresholding, R was normalized using OpenCV's `cv2.normalize()` function. Normalization scales R between 0 and 255, ensuring consistency across different images and lighting conditions.

```

1 # Normalize R for better thresholding
2 R_norm = cv2.normalize(R, None, alpha=0, beta=255, norm_type=cv2.NORM_MINMAX)
3 threshold = threshold_scale * R_norm.max()

```

1.3 Identifying Strong Interest Points

1. Thresholding: We keep only points where the Harris response R exceeds a percentage of its maximum value, ensuring that only strong corners are retained.

2. Local Maxima Search in a 7×7 Window: This ensures that detected keypoints are spatially distinct, preventing clustering of detections around a single corner.

```

1 from scipy.ndimage import maximum_filter
2
3 # Identify local maxima in the response map
4 local_max = (R_norm == maximum_filter(R_norm, size=7))

```

3. Keypoint Storage with Orientation Information: Orientation information is crucial for feature matching in further applications.

```

1 import numpy as np
2
3 # Compute gradient magnitude and orientation
4 magnitude = np.sqrt(Ix**2 + Iy**2)
5 orientation = np.arctan2(Iy, Ix) * (180 / np.pi) # Convert radians to degrees
6
7 # Extract keypoints based on threshold and local maxima
8 for y, x in zip(*np.where(local_max & (R_norm > threshold))):
9     keypoints.append(cv2.KeyPoint(float(x), float(y), 1, orientation[y, x]))

```

1.4 Effect of Image Preprocessing on Corner Detection

Original Image For the original image, no pre-smoothing was applied before calculating the derivatives for detecting interest points. This approach worked perfectly, as smoothing would only reduce the function's quality. The same pre-processing is also for the 180-degree rotated image since the Harris detector is invariant to rotation.

Noisy Image For the noisy image, a Gaussian blur with a 3×3 window was applied before computing derivatives. This preprocessing step helped reduce false detections caused by densely distributed noise pixels. A Gaussian sigma value of 3 was used to ensure effective noise suppression.

```

1 # Apply Gaussian blur to reduce noise
2 gray = cv2.GaussianBlur(gray, (3, 3), 3)

```

Choosing $\sigma = 3$ ensures that only nearby pixels contribute to smoothing, preventing interference from distant noise points.

Pixelated Image The same preprocessing technique as the noisy image was used for the pixelated image. Applying a Gaussian blur helped to smooth out artificial edges and improve the accuracy of derivative calculations.

Bright Image Corner detection on the bright image performed well without additional smoothing, similar to the original image. While the same threshold scaling could be used, a slightly lower threshold was preferred to improve corner detection accuracy.

Blurred Image A major issue with blurred images arises from strong Gaussian smoothing, which significantly reduces image gradients. This results in a more uniform intensity distribution, leading to weak corner responses. The second-moment matrix, which relies on intensity differences, produces small corner responses, causing excessive detections or failing to identify strong features.

To improve corner detection in blurred images, sharpening was applied before computing the gradients. This enhanced edge details, making corner features more distinguishable.

Steps to Improve Detection in Blurred Images: 1. Apply an unsharp masking filter (sharpening) before computing the gradients. 2. Modify Gaussian smoothing parameters to balance detail preservation and noise reduction. 3. Dynamically adjust the threshold based on the image's contrast.

```
1 # Apply sharpening to enhance edges
2 sharpening_kernel = np.array([[[-1, -1, -1],
3                               [-1, 9, -1],
4                               [-1, -1, -1]]])
5 gray = cv2.filter2D(gray, -1, sharpening_kernel)
```

This sharpening step enhances gradients, allowing the detector to identify corners more effectively.

2 Effect of Threshold Value on the Number of Detected Keypoints

Thresholding plays a crucial role in Harris Corner Detection, as it directly impacts the number of keypoints identified. A higher threshold selects only the strongest corners, reducing the total number of detected keypoints, whereas a lower threshold retains weaker corners, potentially increasing noise.

To analyze this effect, I experimented with varying threshold values and plotted the corresponding number of detected keypoints. The threshold values were chosen specifically for each image, as I observed that only within a certain range did significant changes occur. Notably, small changes in the threshold could cause a sharp jump from around 400 detected matches to over 10,000, after which the number stabilized.

My Harris corner detection function employs an adaptive threshold based on the maximum response value R , meaning the threshold operates on a relative scale. Consequently, I input different threshold scale values and also plotted the actual computed threshold values in the graphs. This explains why the numerical values might appear unconventional.

These graphs illustrate how increasing the threshold results in fewer detected keypoints, emphasizing the trade-off between precision and recall in feature detection.

2.1 Visualization of Keypoint Detection at Different Thresholds

Below are visualizations of the number of detected keypoints at various threshold values. Each image demonstrates how increasing the threshold impacts the number of detected corners.

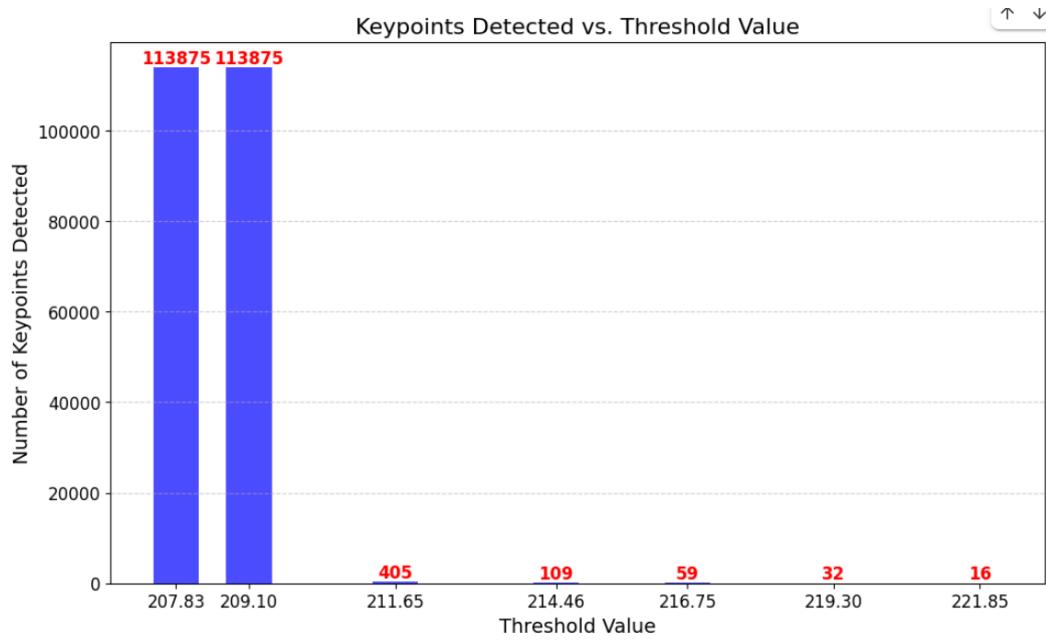


Figure 1: Detected Keypoints in Original Image

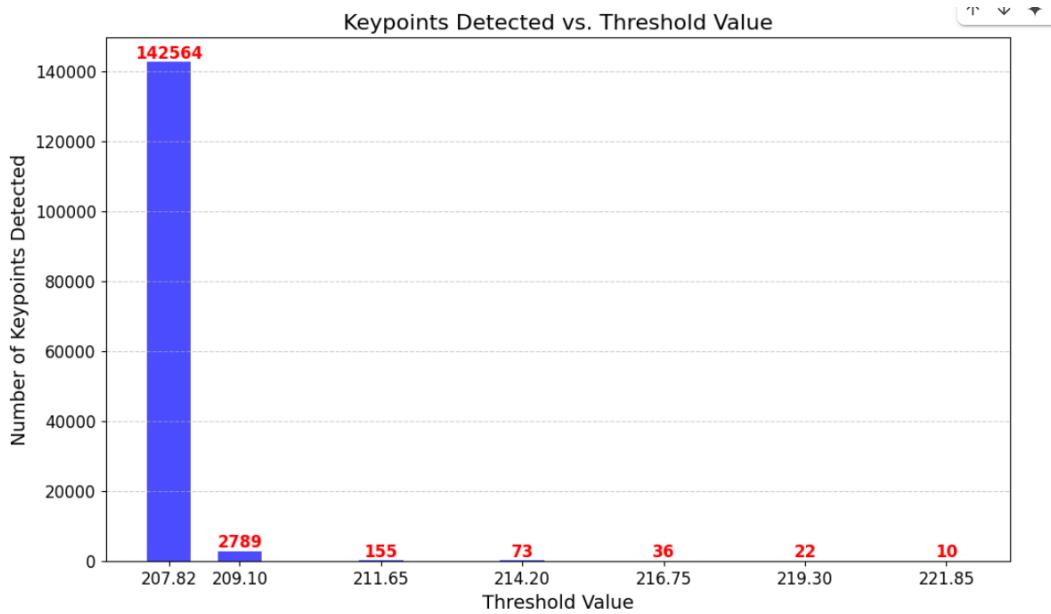


Figure 2: Detected Keypoints in Darker Image

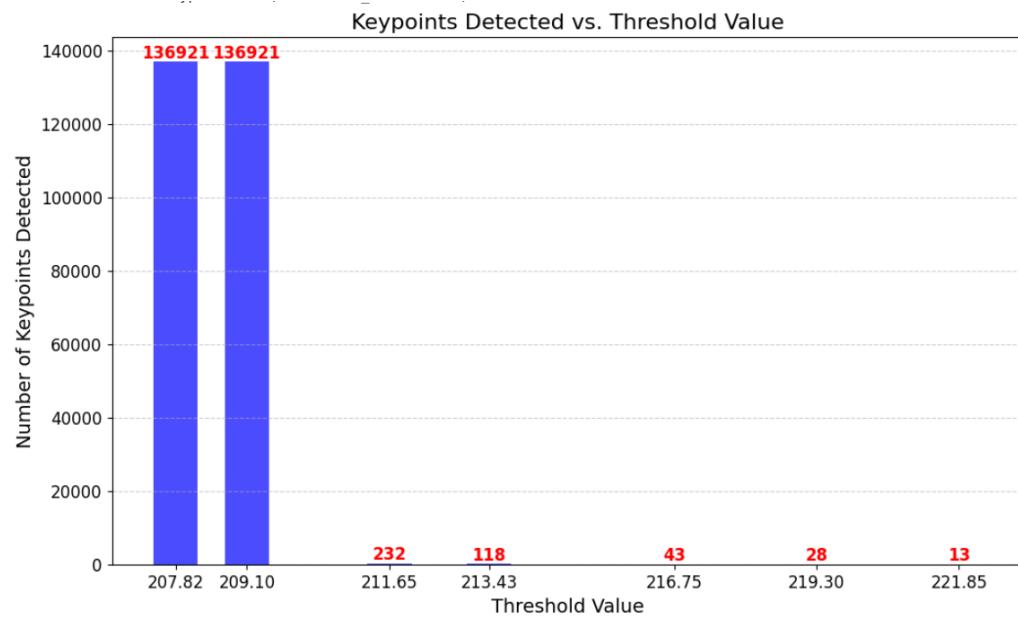


Figure 3: Detected Keypoints in Brighter Image

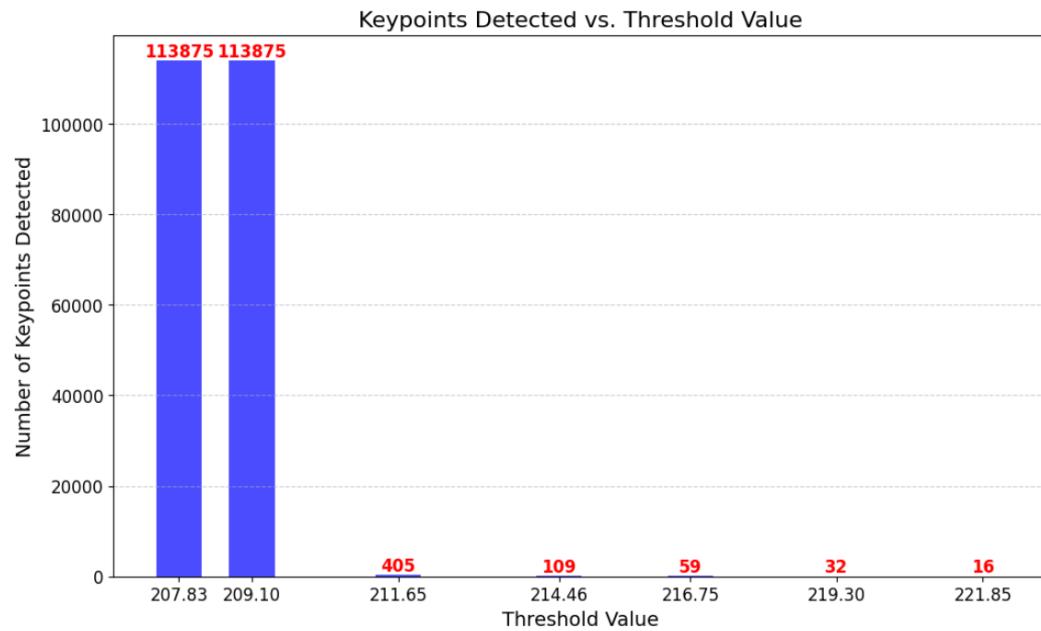


Figure 4: Detected Keypoints in 180-degree Rotated Image

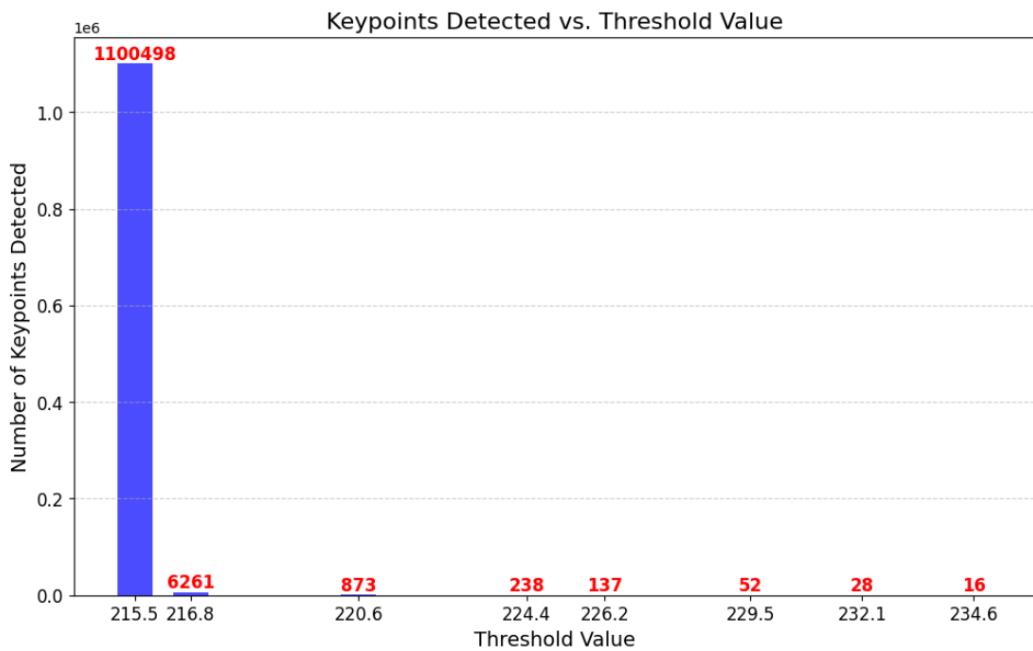


Figure 5: Detected Keypoints in Blurred Image

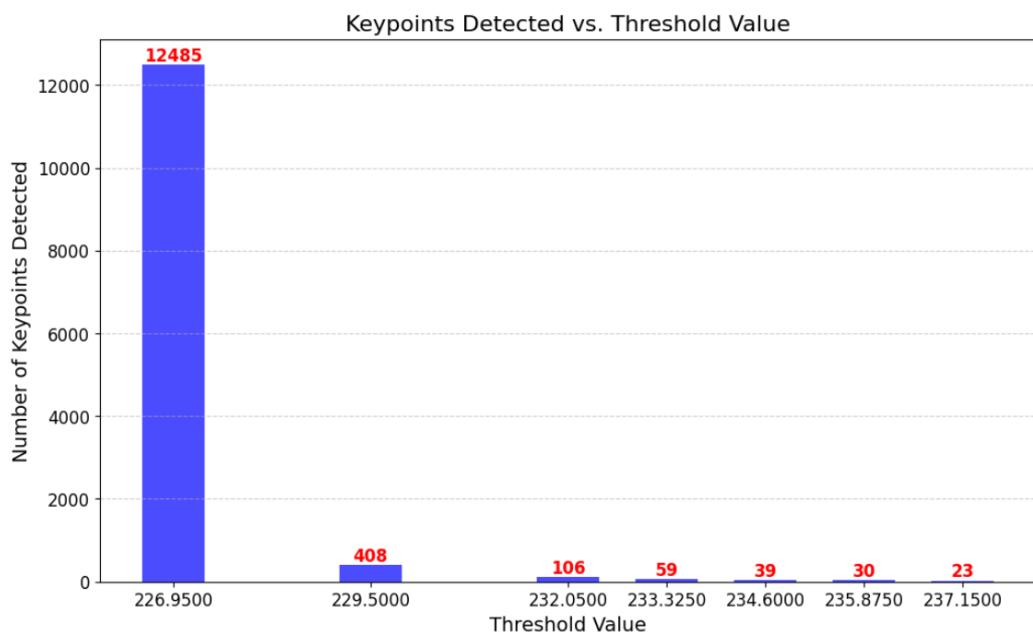


Figure 6: Detected Keypoints in Noisy Image

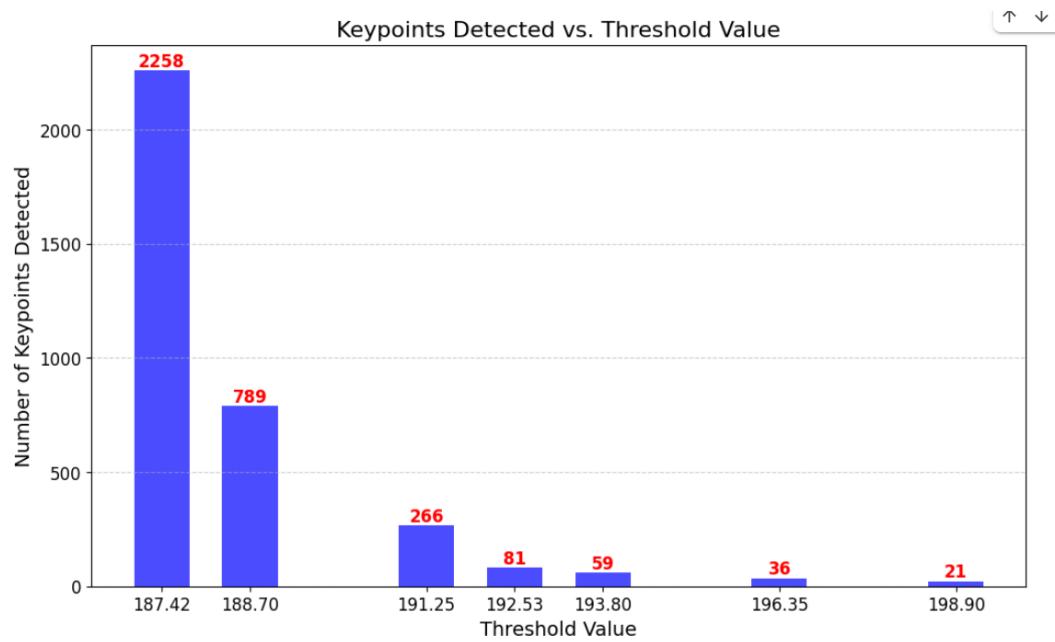


Figure 7: Detected Keypoints in Pixelated Image

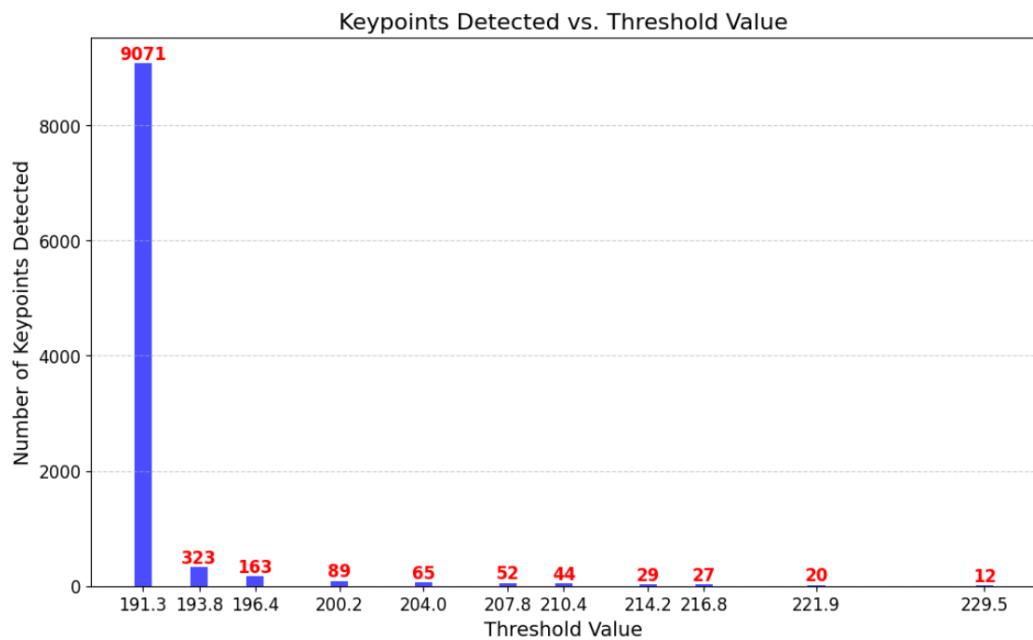


Figure 8: Detected Keypoints in Friends Image

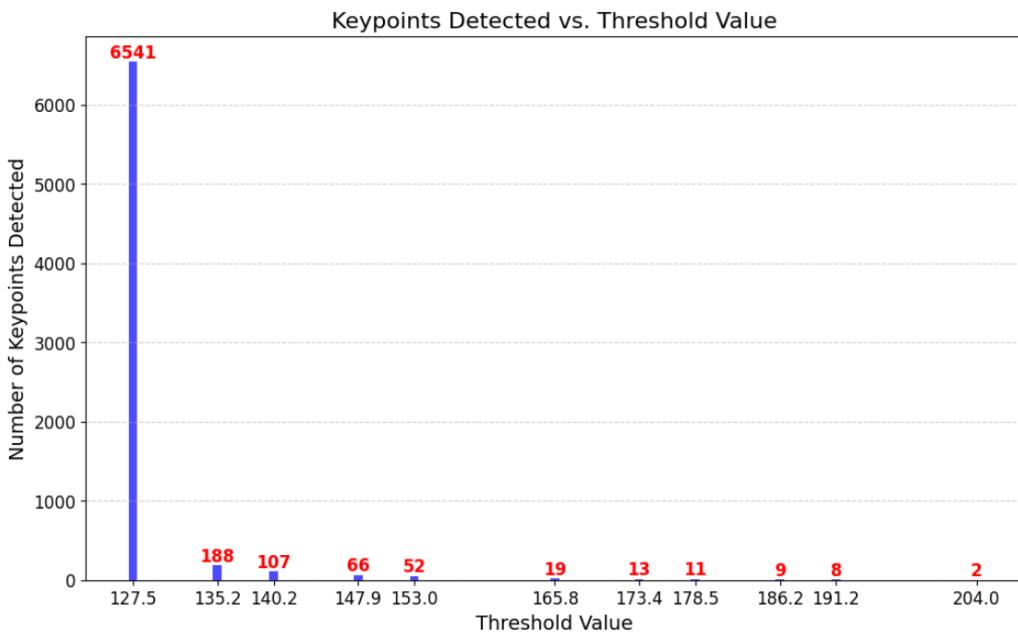


Figure 9: Detected Keypoints in Salon Image

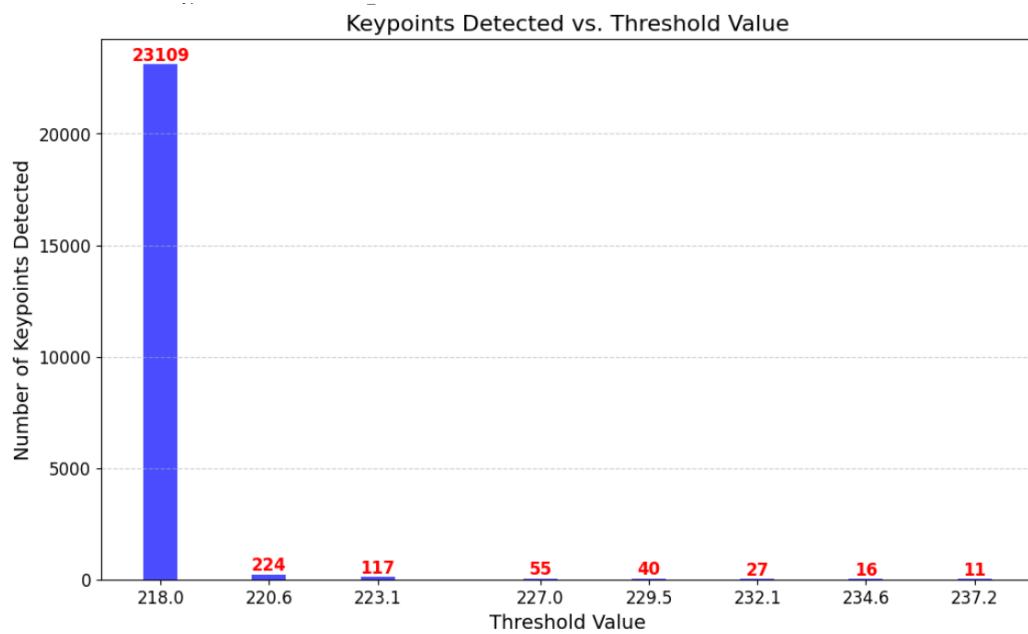


Figure 10: Detected Keypoints in School Image

2 Comparison of ORB Descriptors and Keypoint Detection

To analyze the performance of ORB descriptors, we computed local features using ORB's `compute` function with keypoints detected from our custom Harris detector. Additionally, as an experiment, we utilized the ORB `detect` function to automatically detect keypoints and then applied OpenCV's built-in Harris and FAST methods for further comparison.

The images below showcase keypoints detected using our descriptor with ORB `compute`, as well as keypoints obtained using ORB `detect` with OpenCV Harris and FAST. The images displaying our descriptor include the threshold values used for detection.

2.1 Threshold Values for Different Image Conditions

The following table summarizes the threshold values used for Harris and FAST detectors across different image conditions.

Image Condition	Harris Threshold Scale	FAST Threshold
Original Image	0.01	30
Darker	0.01	30
Brighter	0.01	30
Rotated	0.01	30
Friends (noisy background)	0.05	40
Salon (noisy background)	0.05	40
Blurred	0.08	2
Noisy	0.45	400
Pixelated	0.25	500
School	0.08	50

Table 1: Threshold values for different image conditions

2.2 Analysis of Threshold Selection

The selection of threshold values was influenced by the nature of the images:

- Blurred Images: The threshold for FAST was set to only 2 because blurring reduces the contrast at edges, making keypoints harder to detect. A lower threshold ensures that even faint corners are identified.
- Noisy Images: Higher thresholds were required to filter out false keypoints caused by noise. Harris needed 0.45, while FAST required 400, preventing excessive detections in random noise.
- Pixelated Images: Due to sharp artificial edges in pixelated images, high thresholds (0.25 for Harris and 500 for FAST) helped avoid excessive detection at blocky transitions.
- Noisy Backgrounds (Friends and Salon): A moderate increase in threshold helped reduce the number of non-distinctive keypoints in cluttered regions.
- Brightness and Rotation Variations: The same threshold values were retained for these conditions as they do not significantly affect keypoint localization.

Comparing my Harris implementation with OpenCV's Harris and FAST detectors, FAST produced the best results. This is because FAST is specifically designed for speed and efficiency, making it more robust in detecting keypoints across different transformations. While both Harris versions performed well, especially for rotation and brightness changes, they struggled with noise, blur, and background variations. FAST, on the other hand, adapted better to these challenging conditions due to its corner detection approach, which is less sensitive to such distortions.

```
Number of detected keypoints: 109, threshold_scale: 0.841, threshold: 214.455
```



Figure 1: My Harris + ORB.compute (origg_points.png)

Harris Keypoints - Image 1



Figure 2: ORB.detect + OpenCv Harris (haris_nivno_origg.png)



Figure 3: ORB.detect + OpenCv FAST (fast_origg.png)



Figure 4: My Harris + ORB.compute (dark_points.png)



Figure 5: ORB.detect + OpenCv Harris (haris_nivno_Dark.png)



Figure 6: ORB.detect + OpenCv FAST (fast_dark.png)



Figure 7: My Harris + ORB.compute (bright_points.png)

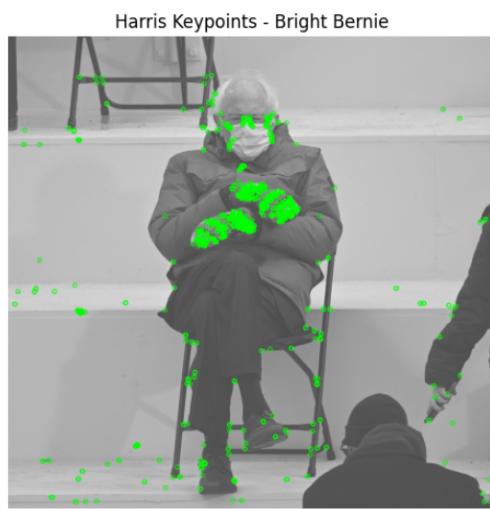


Figure 8: ORB.detect + OpenCv Harris (haris_nivno_bright.png)

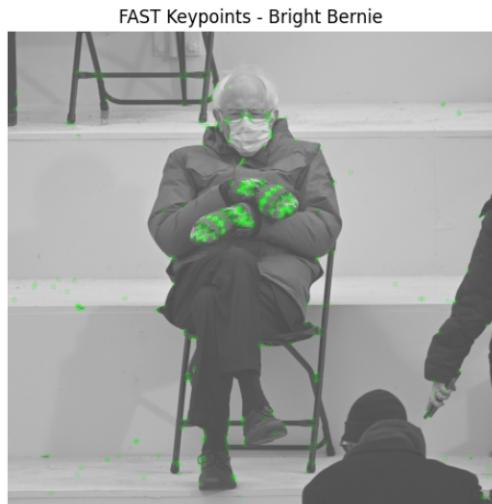


Figure 9: ORB.detect + OpenCv FAST (fast_bright.png)



Figure 10: My Harris + ORB.compute (180-points.png)

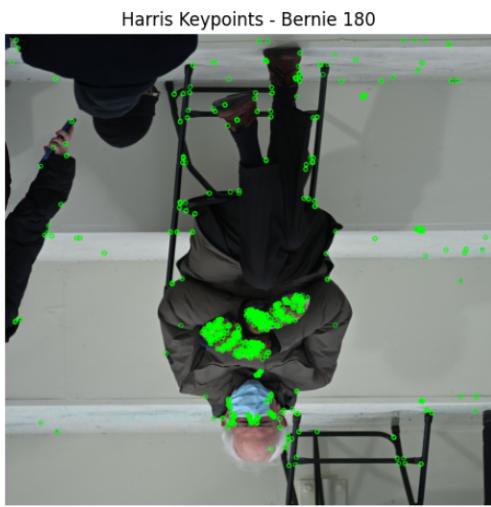


Figure 11: ORB.detect + OpenCv Harris (haris_nivno_180.png)

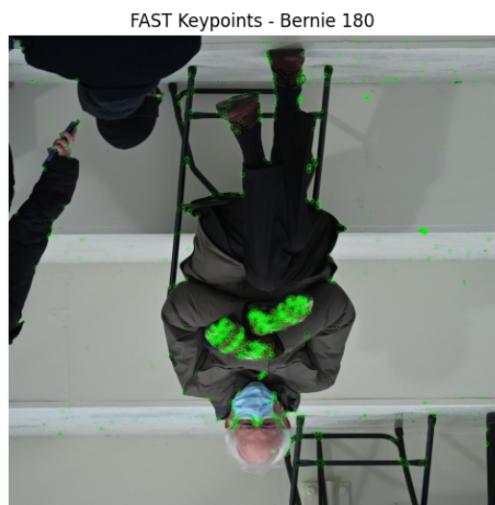


Figure 12: ORB.detect + OpenCv FAST (fast_180.png)

```
Number of detected keypoints: 137, threshold_scale: 0.887, threshold: 226.185
```



Figure 13: My Harris + ORB.compute (blur_points.png)

Harris Keypoints - Blurred



Figure 14: ORB.detect + OpenCv Harris (haris_nivno_blured.png)

FAST Keypoints - Bernie blurred



Figure 15: ORB.detect + OpenCv FAST (fast_blur.png)

Number of detected keypoints: 106, threshold_scale: 0.91, threshold: 232.0499999999998



Figure 16: My Harris + ORB.compute (noisy2_points.png)

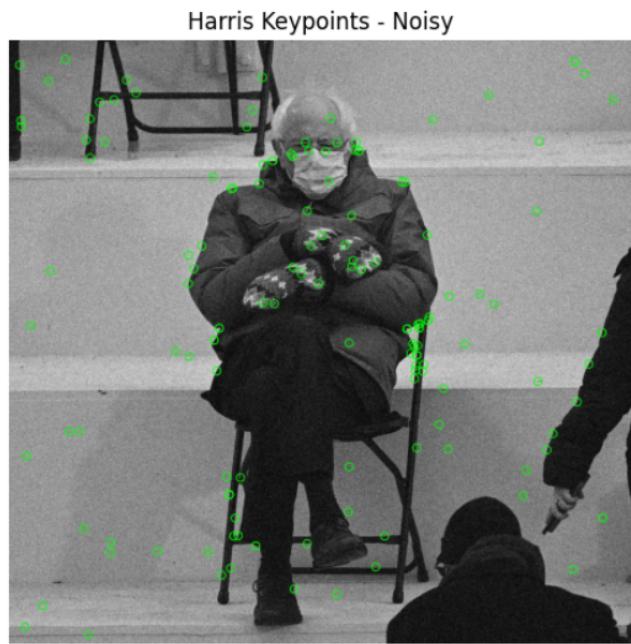


Figure 17: ORB.detect + OpenCv Harris (haris_nivno_noisy.png)



Figure 18: ORB.detect + OpenCv FAST (fats_noisy.png)

```
Number of detected keypoints: 112, threshold_scale: 0.753, threshold: 192.01500000000001
```



Figure 19: My Harris + ORB.compute (pixelated_points.png)

Harris Keypoints - Pixelated



Figure 20: ORB.detect + OpenCv Harris (haris_nivno_pixelated.png)

FAST Keypoints - Pixelated



Figure 21: ORB.detect + OpenCv FAST (fast-pixelated.png)

Number of detected keypoints: 163, threshold scale: 0.77, threshold: 196.35000000000002



Figure 22: My Harris + ORB.compute (frineds-points.png)

Harris Keypoints - Bernie Friends



Figure 23: ORB.detect + OpenCv Harris (haris_nivno_friends.png)

FAST Keypoints - Bernie Friends



Figure 24: ORB.detect + OpenCv FAST (fast_friends.png)

Number of detected keypoints: 107, threshold_scale: 0.55, threshold: 140.25



Figure 25: My Harris + ORB.compute (salon_points.png)

Harris Keypoints - Bright Salon



Figure 26: ORB.detect + OpenCv Harris (haris_nivno_salon.png)

FAST Keypoints - Bright Salon



Figure 27: ORB.detect + OpenCv FAST (fast_salom.png)

Number of detected keypoints: 156, threshold_scale: 0.87, threshold: 221.85



Figure 28: My Harris + ORB.compute (school_points.png)

Harris Keypoints - School



Figure 29: ORB.detect + OpenCv Harris (haris_nivno_school.png)

FAST Keypoints - School



Figure 30: ORB.detect + OpenCv FAST (fast_school.png)

2 Distance functions

As shown in the appendix, I implemented a custom feature matching function using two different approaches: Sum of Squared Differences (SSD) and Lowe's Ratio Test. This function calculates the distance between feature descriptors to identify the best matches between keypoints in two images.

The SSD method computes the squared Euclidean distance between descriptors and selects the closest match for each feature. To discard poor matches, it applies a distance threshold, ensuring only high-confidence matches are retained.

The Ratio Test, on the other hand, helps eliminate ambiguous matches by comparing the best and second-best distances for each feature. If the best match is significantly closer (i.e., the distance is less than a certain ratio of the second-best match), it is accepted; otherwise, it is discarded. This ensures that only distinct, well-matching keypoints are selected, reducing false positives caused by repetitive or noisy features.

This refined approach enhances the accuracy of feature matching, making it more robust against mismatches and improving object detection performance.

3 Matching

Using both the SSD and ratio matching strategies, I obtained the following results. I will present some of them here. For the original image and its 180-degree rotated version, the matching was nearly perfect, as the Harris detector is robust to rotation. The same applied to the darker and brighter images.

However, for the blurred, noisy, and pixelated images, matching became more challenging. This is because the feature calculation window captures noise, and in blurred images, the noise is further smoothed, altering the computed values. Similarly, sharpening affects the gradients, making them slightly different from the original.

The most difficult cases were images with different backgrounds. The descriptors included background information, which differed from the plain white background of the original image. However, I observed that the edges of the gloves were correctly matched, suggesting that even though these images were significantly smaller and their descriptors were calculated at a different scale, the Harris detector was still powerful enough to match them—especially since they shared the same background (Bernie's jacket).

Below, I present some examples of my Harris-based approach using both SSD and ratio matching. Additionally, I include one example using OpenCV's Harris and FAST detectors, which produced slightly stronger results compared to mine.

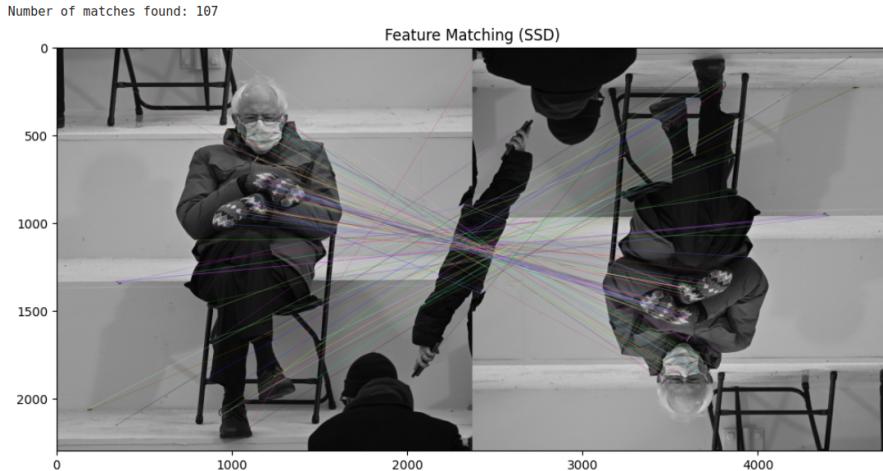


Figure 1: My Haris: 180 - SSD

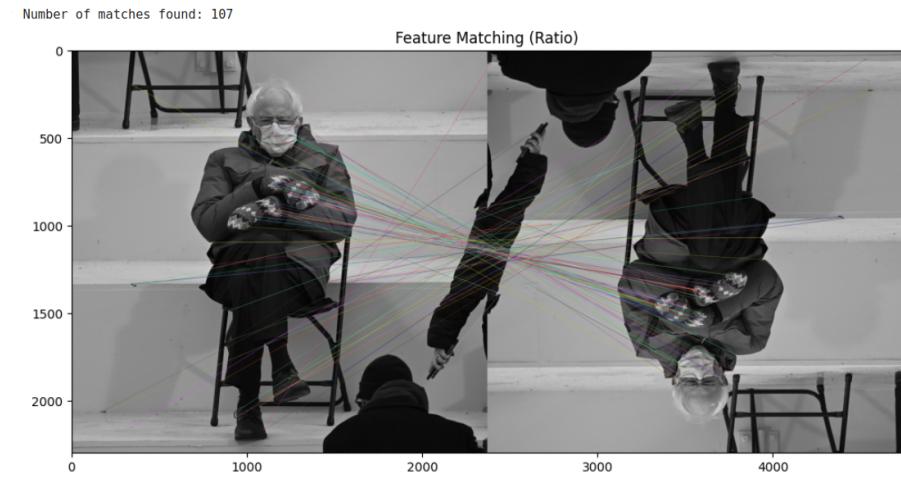


Figure 2: My Haris: 180 - Ratio

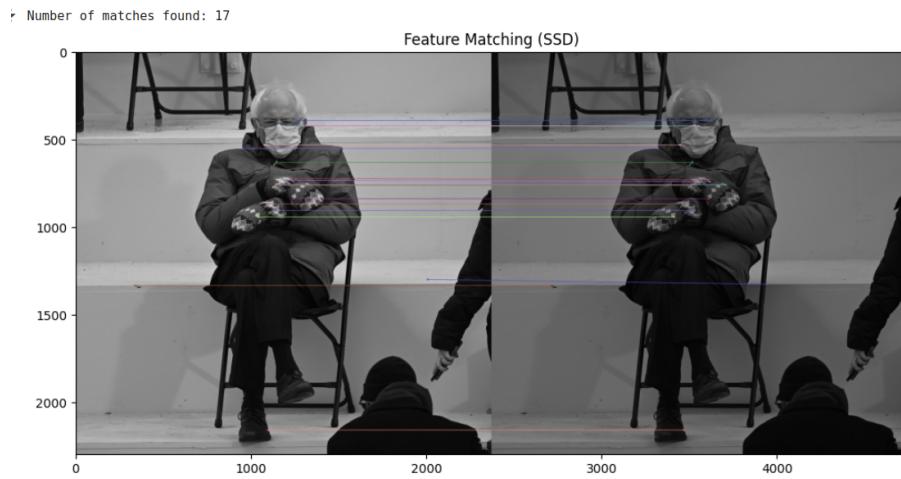


Figure 3: My Haris: Dark - SSD

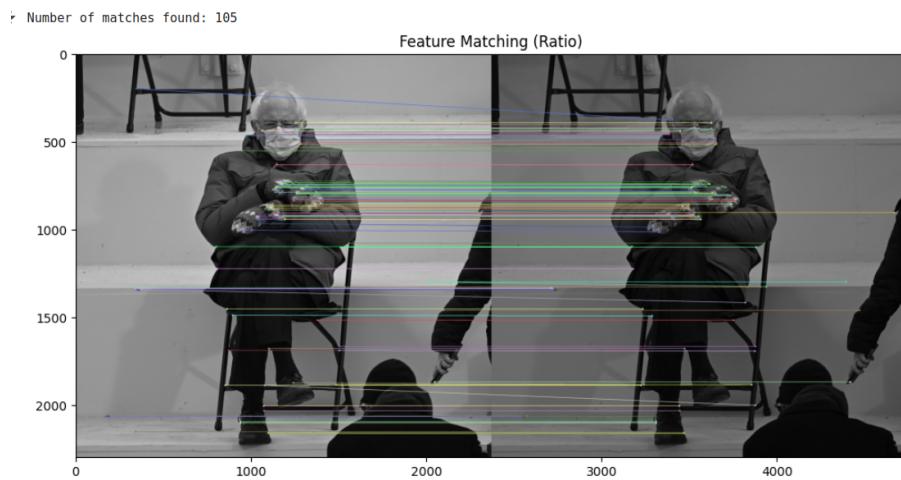


Figure 4: My Haris: Dark - Ratio

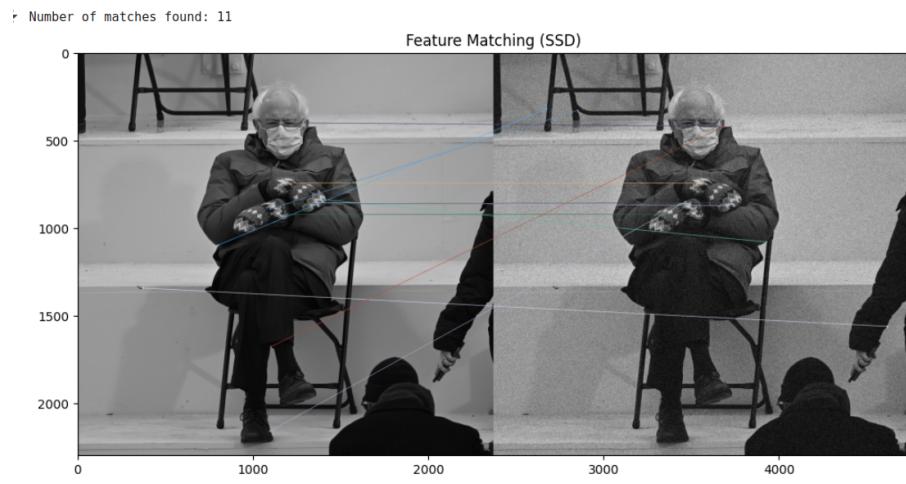


Figure 5: My Haris: Noisy - SSD

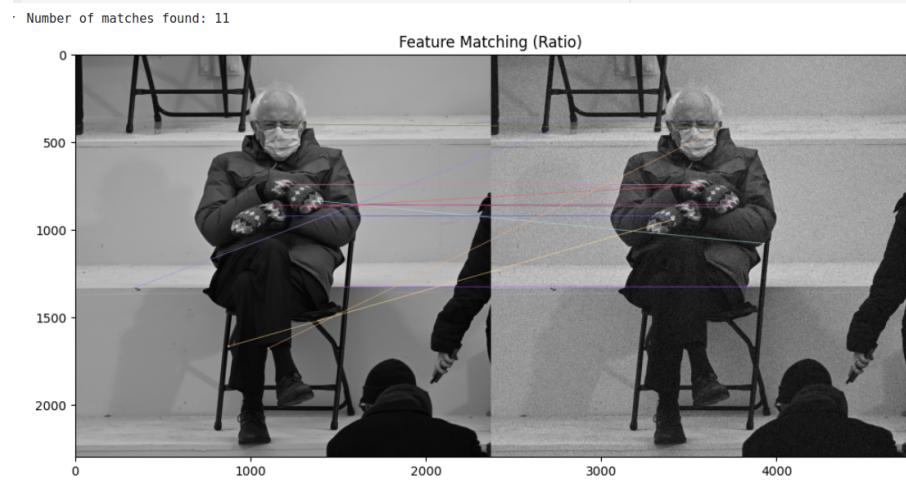


Figure 6: My Haris: Noisy - Ratio

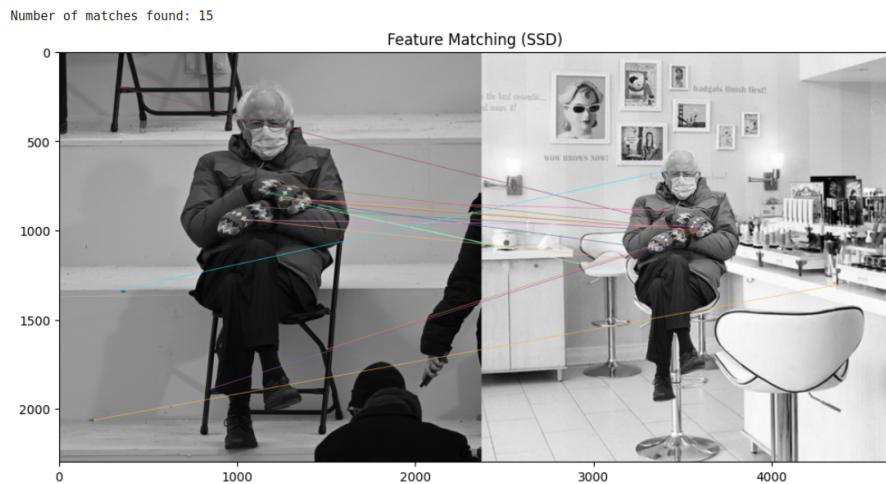


Figure 7: My Haris: Salon - SSD

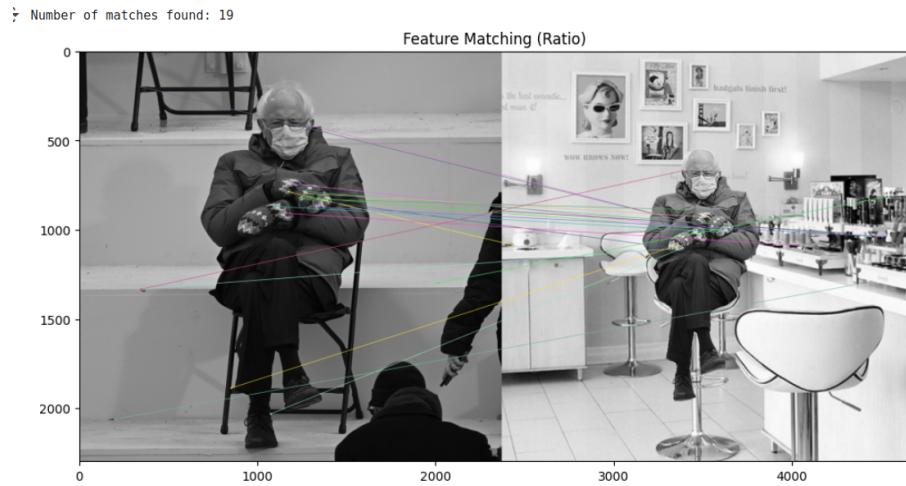


Figure 8: My Haris: Salon - Ratio

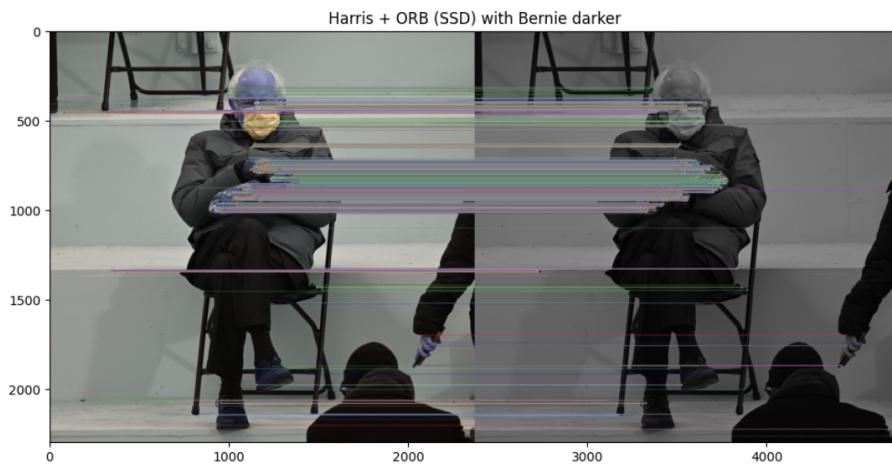


Figure 9: Their Haris: Dark - SSD

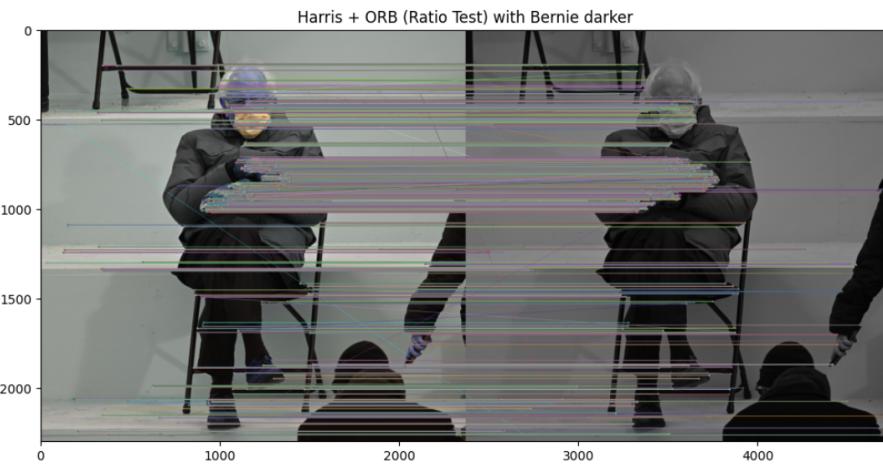


Figure 10: Their Haris: Dark - Ratio

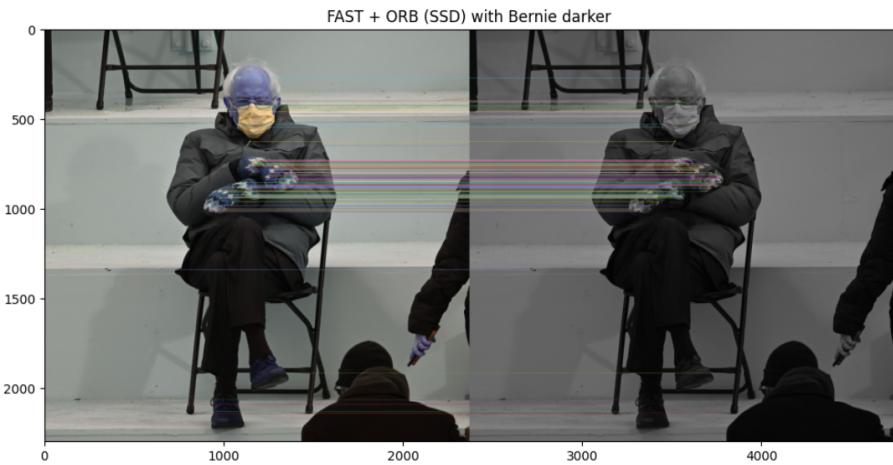


Figure 11: Their FAST: Dark - SSD

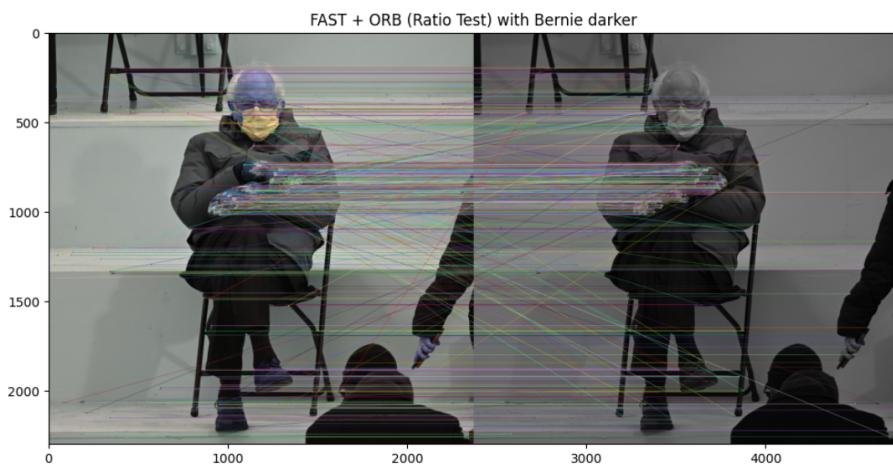


Figure 12: Their FAST: Dark - Ratio

A Appendix: Code Implementation

```
1 import cv2
2 import numpy as np
3 import scipy.ndimage
4 from google.colab.patches import cv2_imshow
5
6 def HarrisPointsDetector(image, k=0.05, threshold_scale=0.8, visualize=True):
7     """
8         Detects keypoints using the Harris Corner Detection method.
9
10    Parameters:
11        image (numpy.ndarray): Input image (BGR or grayscale).
12        k (float): Harris corner constant.
13        threshold_scale (float): Scale for selecting strong keypoints.
14        visualize (bool): Whether to display detected corners.
15
16    Returns:
17        keypoints (list of cv2.KeyPoint): List of OpenCV KeyPoint objects.
18    """
19
20    # Convert to grayscale if needed
21    if len(image.shape) == 3:
22        gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
23    else:
24        gray = image.copy()
25
26    gray = np.float32(gray) # Convert to float32 for precision
27
28    #gray = cv2.GaussianBlur(gray, (3, 3), 3)
29
30    # #Apply sharpening to counter excessive blurring
31    # sharpening_kernel = np.array([[-1, -1, -1],
32    #                               [-1, 9, -1],
33    #                               [-1, -1, -1]])
34    # gray = cv2.filter2D(gray, -1, sharpening_kernel)
35
36
37
38
39    # Compute gradients using 3x3 Sobel operator (with reflection padding)
40    Ix = cv2.Sobel(gray, cv2.CV_64F, 1, 0, ksize=3, borderType=cv2.BORDER_REFLECT)
41    Iy = cv2.Sobel(gray, cv2.CV_64F, 0, 1, ksize=3, borderType=cv2.BORDER_REFLECT)
42
43    # Compute gradient magnitude and orientation (in degrees)
44    magnitude = np.sqrt(Ix**2 + Iy**2)
45    orientation = np.arctan2(Iy, Ix) * (180 / np.pi) # Convert radians to degrees
46
47    # Compute second-moment matrix elements
48    Ixx = Ix ** 2
49    Iyy = Iy ** 2
50    Ixy = Ix * Iy
51
52    # Apply 5x5 Gaussian weighting (sigma = 0.5)
53    gaussian_filter = (5, 5)
54    sigma = 0.5
55    Sxx = cv2.GaussianBlur(Ixx, gaussian_filter, sigma)
56    Syy = cv2.GaussianBlur(Iyy, gaussian_filter, sigma)
57    Sxy = cv2.GaussianBlur(Ixy, gaussian_filter, sigma)
58
59    # Compute Harris response
60    detM = (Sxx * Syy) - (Sxy ** 2)
61    traceM = Sxx + Syy
62    R = detM - k * (traceM ** 2)
63
64    # Normalize R for thresholding
65    R_norm = cv2.normalize(R, None, alpha=0, beta=255, norm_type=cv2.NORM_MINMAX)
66    threshold = threshold_scale * R_norm.max()
67
68    # Find local maxima in a 7x7 neighborhood
69    local_max = (R_norm == scipy.ndimage.maximum_filter(R_norm, size=7))
```

```

70
71     keypoints = []
72     for y, x in zip(*np.where(local_max & (R_norm > threshold))):
73         keypoints.append(cv2.KeyPoint(float(x), float(y), 1, orientation[y, x])) # Store
74             orientation
75
76     print(f"Number of detected keypoints: {len(keypoints)}, threshold_scale: {threshold_scale}, threshold: {threshold}")
77
78     # Visualization
79     if visualize:
80         output_image = cv2.cvtColor(gray, cv2.COLOR_GRAY2BGR)
81         for kp in keypoints:
82             cv2.circle(output_image, (int(kp.pt[0]), int(kp.pt[1])), 4, (0, 0, 255), -1) # Red
83                 dots for keypoints
84
85     # Resize the output image (to 50% of its original size)
86     output_image_resized = cv2.resize(output_image, (0, 0), fx=0.5, fy=0.5) # Resize by 50%
87     cv2_imshow(output_image_resized)
88
89     return keypoints, threshold
90
91
92 image = cv2.imread("bernieSchoolLunch.jpeg")
93
94 # Detect keypoints
95 keypoints = HarrisPointsDetector(image, k=0.05, threshold_scale=0.87)
96
97
98 import cv2
99 import numpy as np
100 import scipy.ndimage
101 import matplotlib.pyplot as plt
102
103 def evaluate_thresholds(image):
104     threshold_scales = [0.855, 0.865, 0.875, 0.89, 0.9, 0.91, 0.92, 0.93]
105     num_keypoints = []
106     actual_thresholds = []
107
108     for threshold_scale in threshold_scales:
109         keypoints, threshold = HarrisPointsDetector(image, threshold_scale=threshold_scale,
110             visualize=False)
111         num_keypoints.append(len(keypoints))
112         actual_thresholds.append(threshold)
113
114     plt.figure(figsize=(12, 7))
115     bars = plt.bar(actual_thresholds, num_keypoints, color='b', alpha=0.7)
116
117     plt.xticks(actual_thresholds, labels=[f"{t:.1f}" for t in actual_thresholds], fontsize=12)
118     plt.yticks(fontsize=12)
119     plt.xlabel("Threshold Value", fontsize=14)
120     plt.ylabel("Number of Keypoints Detected", fontsize=14)
121     plt.title("Keypoints Detected vs Threshold Value", fontsize=16)
122     plt.grid(axis='y', linestyle='--', alpha=0.6)
123
124     for bar, txt in zip(bars, num_keypoints):
125         plt.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 0.5, str(txt),
126             ha='center', va='bottom', fontsize=12, color='red', weight='bold')
127
128     plt.show()
129
130
131 image = cv2.imread("bernieSchoolLunch.jpeg")
132 evaluate_thresholds(image)
133
134 image_origg = cv2.imread("bernieSanders.jpg")
135
136 # Detect keypoints
137 keypoints_origg, threshold_origg = HarrisPointsDetector(image_origg, threshold_scale=0.841)
138
139 image_noisy = cv2.imread("bernieNoisy2.png")

```

```

137
138 # Detect keypoints
139 keypoints_noisy, threshold_noisy = HarrisPointsDetector(image_noisy, threshold_scale=0.91)
140
141 image_friends = cv2.imread("BernieFriends.png")
142
143 # Detect keypoints
144 keypoints_friends, _ = HarrisPointsDetector(image_friends, threshold_scale=0.77)
145
146 image_school = cv2.imread("bernieShoolLunch.jpeg")
147
148 # Detect keypoints
149 keypoints_school, _ = HarrisPointsDetector(image_school, threshold_scale=0.87)
150
151
152 image.blur = cv2.imread("bernieMoreblurred.jpg")
153
154 # Detect keypoints
155 keypoints.blur, _ = HarrisPointsDetector(image.blur, threshold_scale=0.887)
156
157 image.180 = cv2.imread("bernie180.jpg")
158
159 # Detect keypoints
160 keypoints.180, _ = HarrisPointsDetector(image.180, threshold_scale=0.841)
161
162 image.pixel = cv2.imread("berniePixelated2.png")
163
164 # Detect keypoints
165 keypoints.pixel, _ = HarrisPointsDetector(image.pixel, threshold_scale=0.75)
166
167 image.bright = cv2.imread("brighterBernie.jpg")
168
169 # Detect keypoints
170 keypoints.bright, _ = HarrisPointsDetector(image.bright, threshold_scale=0.837)
171
172
173 image.dark = cv2.imread("darkerBernie.jpg")
174
175 # Detect keypoints
176 keypoints.dark, _ = HarrisPointsDetector(image.dark, threshold_scale=0.83)
177
178 image.salon = cv2.imread("bernieBenefitBeautySalon.jpeg")
179
180 # Detect keypoints
181 keypoints.salon, _ = HarrisPointsDetector(image.salon, threshold_scale=0.55)
182
183 import cv2
184 import numpy as np
185
186 def featureDescriptor(image, keypoints):
187     """
188         Computes feature descriptors for given keypoints using ORB.
189
190     Parameters:
191         image (numpy.ndarray): Input grayscale image.
192         keypoints (list of cv2.KeyPoint): Detected keypoints with location and orientation.
193
194     Returns:
195         descriptors (numpy.ndarray): ORB descriptors of shape (num_keypoints, feature_dim).
196     """
197
198     # Convert image to grayscale
199     if len(image.shape) == 3:
200         gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
201     else:
202         gray = image.copy()
203
204     # Create ORB detector
205     orb = cv2.ORB_create(500)
206
207

```

```

208     if not keypoints:
209         print("No keypoints detected!")
210         return [], np.array([])
211
212
213
214     # Compute descriptors using ORB
215     keypoints, descriptors = orb.compute(gray, keypoints)
216
217     if descriptors is None:
218         print("No descriptors found!")
219         return np.array([]) # Return empty array if no descriptors are found
220
221     print(f"Computed {descriptors.shape[0]} descriptors with dimension {descriptors.shape[1]}")
222
223     return keypoints, descriptors
224
225
226 image_origg = cv2.imread("bernieSanders.jpg")
227
228 _, descriptors_origg = featureDescriptor(image_origg, keypoints_origg)
229
230 # Print descriptor shape
231 print(f"Descriptor shape: {descriptors_origg.shape}")
232 image_school = cv2.imread("bernieShoolLunch.jpeg")
233 # Step 2: Compute ORB descriptors for these keypoints
234 _, descriptors_school = featureDescriptor(image_school, keypoints_school)
235
236
237 print(f"Descriptor shape: {descriptors_school.shape}")
238
239 image_friends = cv2.imread("BernieFriends.png")
240 # Step 2: Compute ORB descriptors for these keypoints
241 _, descriptors_friends = featureDescriptor(image_friends, keypoints_friends)
242
243 # Print descriptor shape
244 print(f"Descriptor shape: {descriptors_friends.shape}")
245 image_noisy = cv2.imread("bernieNoisy2.png")
246 # Step 2: Compute ORB descriptors for these keypoints
247 _, descriptors_noisy = featureDescriptor(image_noisy, keypoints_noisy)
248
249 # Print descriptor shape
250 print(f"Descriptor shape: {descriptors_noisy.shape}")
251
252 image_180 = cv2.imread("bernie180.jpg")
253 # Step 2: Compute ORB descriptors for these keypoints
254 _, descriptors_180 = featureDescriptor(image_180, keypoints_180)
255
256 # Print descriptor shape
257 print(f"Descriptor shape: {descriptors_180.shape}")
258
259 image_bright = cv2.imread("brighterBernie.jpg")
260 # Step 2: Compute ORB descriptors for these keypoints
261 _, descriptors_bright = featureDescriptor(image_bright, keypoints_bright)
262
263 # Print descriptor shape
264 print(f"Descriptor shape: {descriptors_bright.shape}")
265
266 image_dark = cv2.imread("darkerBernie.jpg")
267 # Step 2: Compute ORB descriptors for these keypoints
268 _, descriptors_dark = featureDescriptor(image_dark, keypoints_dark)
269
270 # Print descriptor shape
271 print(f"Descriptor shape: {descriptors_dark.shape}")
272
273 image_pixel = cv2.imread("berniePixelated2.png")
274 # Step 2: Compute ORB descriptors for these keypoints
275 _, descriptors_pixel = featureDescriptor(image_pixel, keypoints_pixel)
276
277 # Print descriptor shape
278 print(f"Descriptor shape: {descriptors_pixel.shape}")

```

```

279
280     image_school = cv2.imread("bernieShoolLunch.jpeg")
281     # Step 2: Compute ORB descriptors for these keypoints
282     _, descriptors_school = featureDescriptor(image_school, keypoints_school)
283
284     # Print descriptor shape
285     print(f"Descriptor shape: {descriptors_school.shape}")
286
287
288     image.blur = cv2.imread("bernieMoreblurred.jpg")
289     # Step 2: Compute ORB descriptors for these keypoints
290     _, descriptors.blur = featureDescriptor(image.blur, keypoints.blur)
291
292     # Print descriptor shape
293     print(f"Descriptor shape: {descriptors.blur.shape}")
294
295     image.salon = cv2.imread("bernieBenefitBeautySalon.jpeg")
296     # Step 2: Compute ORB descriptors for these keypoints
297     _, descriptors.salon = featureDescriptor(image.salon, keypoints.salon)
298
299     # Print descriptor shape
300     print(f"Descriptor shape: {descriptors.salon.shape}")
301
302
303
304     import numpy as np
305     import cv2
306     import matplotlib.pyplot as plt
307     import random
308     from scipy.spatial.distance import cdist
309
310     def match_and_visualize(image1, keypoints1, descriptors1, image2, keypoints2, descriptors2,
311                             matcher_type="Ratio", max_distance=85000, ratio_threshold=0.8):
312         """
313             Matches features between two images and visualizes them with vibrant colored lines.
314
315             Parameters:
316                 image1, image2: Input images.
317                 keypoints1, keypoints2: Detected keypoints.
318                 descriptors1, descriptors2: Feature descriptors.
319                 matcher_type: "SSD" for Sum of Squared Differences, "Ratio" for Lowe's Ratio Test.
320                 max_distance: Threshold for SSD distance filtering.
321                 ratio_threshold: Ratio threshold for Lowe's Ratio Test.
322
323             Returns:
324                 Matches visualization.
325
326             # Resize second image to match height of first image
327             h1, w1 = image1.shape[:2]
328             h2, w2 = image2.shape[:2]
329             scale_factor = h1 / h2
330             image2_resized = cv2.resize(image2, (int(w2 * scale_factor), h1))
331
332             # Convert images to grayscale if needed
333             image1_gray = cv2.cvtColor(image1, cv2.COLOR_BGR2GRAY) if len(image1.shape) == 3 else image1
334             image2_gray = cv2.cvtColor(image2_resized, cv2.COLOR_BGR2GRAY) if len(image2.shape) == 3 else
335                         image2_resized
336
337             # Ensure descriptors exist
338             if descriptors1 is None or descriptors2 is None:
339                 print("No descriptors found in one or both images!")
340                 return []
341
342             # Convert descriptors to float for compatibility with cdist
343             descriptors1 = descriptors1.astype(np.float32)
344             descriptors2 = descriptors2.astype(np.float32)
345
346             matches = []
347             if matcher_type == "SSD":

```

```

348     # Compute SSD matrix
349     ssd_matrix = cdist(descriptors1, descriptors2, metric='sqeuclidean')
350     best_matches = np.argmin(ssd_matrix, axis=1)
351
352     for i, j in enumerate(best_matches):
353         distance = ssd_matrix[i, j]
354         if distance < max_distance: # Thresholding to remove bad matches
355             matches.append(cv2.DMatch(_queryIdx=i, _trainIdx=j, _distance=distance))
356
357 elif matcher_type == "Ratio":
358     # Compute SSD and apply Ratio Test
359     ssd_matrix = cdist(descriptors1, descriptors2, metric='sqeuclidean')
360     sorted_indices = np.argsort(ssd_matrix, axis=1)[:, :2] # Get top 2 matches
361
362     for i in range(len(descriptors1)):
363         best_idx, second_best_idx = sorted_indices[i]
364         best_distance, second_best_distance = ssd_matrix[i, best_idx], ssd_matrix[i,
365             second_best_idx]
366
367         # Apply Lowe's Ratio Test
368         if best_distance < ratio_threshold * second_best_distance:
369             matches.append(cv2.DMatch(_queryIdx=i, _trainIdx=best_idx, _distance=best_distance
370             ))
371
372     print(f"Number of matches found: {len(matches)}")
373
374     # Create a combined image for drawing matches
375     output_image = np.zeros((h1, w1 + image2_resized.shape[1], 3), dtype=np.uint8)
376
377     # Convert grayscale images to 3-channel for color drawing
378     output_image[:, :, :w1] = cv2.cvtColor(image1_gray, cv2.COLOR_GRAY2BGR)
379     output_image[:, :, w1:] = cv2.cvtColor(image2_gray, cv2.COLOR_GRAY2BGR)
380
381     # Draw matches with vibrant colors
382     for match in matches:
383         pt1 = tuple(map(int, keypoints1[match.queryIdx].pt))
384         pt2 = tuple(map(int, keypoints2[match.trainIdx].pt))
385         pt2 = (int(pt2[0] * scale_factor) + w1, int(pt2[1] * scale_factor)) # Adjust pt2 after
386         resizing
387
388         # Generate a vibrant random color
389         color = tuple(random.randint(50, 255) for _ in range(3))
390
391         # Draw line between matched points
392         cv2.line(output_image, pt1, pt2, color, thickness=2)
393
394         # Draw keypoints with bold circles
395         cv2.circle(output_image, pt1, 5, color, -1)
396         cv2.circle(output_image, pt2, 5, color, -1)
397
398     # Show the results
399     plt.figure(figsize=(12, 6))
400     plt.imshow(cv2.cvtColor(output_image, cv2.COLOR_BGR2RGB))
401     plt.title(f"Feature Matching {matcher_type}")
402     plt.show()
403
404     return matches
405
406
407 from google.colab.patches import cv2_imshow
408
409 match_and_visualize(image_origg, keypoints_origg, descriptors_origg, image_school,
410     keypoints_school, descriptors_school, matcher_type="Ratio")
411
412
413 import cv2
414 import numpy as np
415 import matplotlib.pyplot as plt
416 import random
417 from scipy.spatial.distance import cdist
418
419 def harris_corner_detector(image, block_size=2, ksize=3, k=0.05, threshold=0.08):

```

```

415     """
416     Detects keypoints using Harris Corner Detector.
417
418     Parameters:
419         image: Grayscale input image.
420         block_size: Neighborhood size.
421         ksize: Aperture parameter for the Sobel operator.
422         k: Harris detector free parameter.
423         threshold: Threshold for selecting keypoints.
424
425     Returns:
426         Keypoints list.
427     """
428     harris_response = cv2.cornerHarris(image, block_size, ksize, k)
429     harris_response = cv2.dilate(harris_response, None) # Dilate for better visualization
430     keypoints = np.argwhere(harris_response > threshold * harris_response.max())
431     keypoints = [cv2.KeyPoint(float(x[1]), float(x[0]), 4) for x in keypoints]
432     return keypoints
433
434 def fast_detector(image, threshold=50):
435     """
436     Detects keypoints using FAST feature detector.
437
438     Parameters:
439         image: Grayscale input image.
440         threshold: FAST threshold.
441
442     Returns:
443         Keypoints list.
444     """
445     fast = cv2.FastFeatureDetector_create(threshold)
446     keypoints = fast.detect(image, None)
447     return keypoints
448
449 def compute_orb_descriptors(image, keypoints):
450     """
451     Computes ORB descriptors for the given keypoints.
452
453     Parameters:
454         image: Grayscale input image.
455         keypoints: Keypoints list.
456
457     Returns:
458         Keypoints and descriptors.
459     """
460     orb = cv2.ORB_create()
461     keypoints, descriptors = orb.compute(image, keypoints)
462     return keypoints, descriptors
463
464 def match_features_ssd(descriptors1, descriptors2, max_distance=100):
465     """
466     Matches features using Sum of Squared Differences (SSD).
467
468     Parameters:
469         descriptors1, descriptors2: Feature descriptors.
470         max_distance: Threshold for SSD distance filtering.
471
472     Returns:
473         List of cv2.DMatch objects.
474     """
475     if descriptors1 is None or descriptors2 is None:
476         return []
477
478     descriptors1 = descriptors1.astype(np.float32)
479     descriptors2 = descriptors2.astype(np.float32)
480
481     ssd_matrix = cdist(descriptors1, descriptors2, metric='sqeuclidean')
482     best_matches = np.argmin(ssd_matrix, axis=1)
483
484     matches = []
485     for i, j in enumerate(best_matches):

```

```

486     distance = ssd_matrix[i, j]
487     if distance < max_distance:
488         matches.append(cv2.DMatch(_queryIdx=i, _trainIdx=j, _distance=distance))
489
490     return matches
491
492 def match_features_ratio(descriptors1, descriptors2, ratio_threshold=0.9):
493     """
494     Matches features using Lowe's Ratio Test.
495
496     Parameters:
497         descriptors1, descriptors2: Feature descriptors.
498         ratio_threshold: Ratio test threshold.
499
500     Returns:
501         List of cv2.DMatch objects.
502     """
503     if descriptors1 is None or descriptors2 is None:
504         return []
505
506     descriptors1 = descriptors1.astype(np.float32)
507     descriptors2 = descriptors2.astype(np.float32)
508
509     ssd_matrix = cdist(descriptors1, descriptors2, metric='sqeuclidean')
510     sorted_indices = np.argsort(ssd_matrix, axis=1)[:, :2] # Get top 2 matches
511
512     matches = []
513     for i in range(len(descriptors1)):
514         best_idx, second_best_idx = sorted_indices[i]
515         best_distance, second_best_distance = ssd_matrix[i, best_idx], ssd_matrix[i,
516                                         second_best_idx]
517
518         if best_distance < ratio_threshold * second_best_distance:
519             matches.append(cv2.DMatch(_queryIdx=i, _trainIdx=best_idx, _distance=best_distance))
520
521     return matches
522
523 def visualize_matches(image1, keypoints1, image2, keypoints2, matches, title):
524     """
525     Visualizes feature matches with vibrant colors.
526
527     Parameters:
528         image1, image2: Input images.
529         keypoints1, keypoints2: Keypoints.
530         matches: List of cv2.DMatch objects.
531         title: Title of the plot.
532
533     Returns:
534         None
535     """
536     img_matches = cv2.drawMatches(image1, keypoints1, image2, keypoints2, matches, None, flags=cv2
537                                 .DrawMatchesFlags_NOT_DRAW_SINGLE_POINTS)
538
539     plt.figure(figsize=(12, 6))
540     plt.imshow(img_matches)
541     plt.title(title)
542     plt.show()
543
544 def visualize_keypoints(image, keypoints, title="Keypoints Visualization"):
545     """
546     Draws keypoints on an image and displays it with larger dots.
547
548     Parameters:
549         image: Original image.
550         keypoints: List of detected keypoints.
551         title: Title for the visualization.
552
553     Returns:
554         None
555     """
556     # Increase the size of keypoints by setting a custom diameter

```

```

555     for kp in keypoints:
556         kp.size = 10 # Adjust size as needed
557
558     image_with_keypoints = cv2.drawKeypoints(image, keypoints, None, color=(0, 255, 0), flags=cv2.
559                                             DRAW_MATCHES_FLAGS_DRAW_RICH_KEYPOINTS)
560
561     plt.figure(figsize=(6, 6))
562     plt.imshow(cv2.cvtColor(image_with_keypoints, cv2.COLOR_BGR2RGB))
563     plt.title(title)
564     plt.axis("off")
565     plt.show()
566
567
568
569 # Load images
570 image1 = cv2.imread("bernieShoolLunch.jpeg")
571 image2 = cv2.imread("darkerBernie.jpg")
572
573
574 # Convert to grayscale
575 image1_gray = cv2.cvtColor(image1, cv2.COLOR_BGR2GRAY)
576 image2_gray = cv2.cvtColor(image2, cv2.COLOR_BGR2GRAY)
577
578
579 # Feature detection using Harris
580 harris_keypoints1 = harris_corner_detector(image1_gray)
581 harris_keypoints2 = harris_corner_detector(image2_gray)
582 # Feature detection using FAST
583 fast_keypoints1 = fast_detector(image1_gray)
584 fast_keypoints2 = fast_detector(image2_gray)
585
586
587 # Compute ORB descriptors
588 harris_descriptors1, harris_descriptors2 = compute_orb_descriptors(image1_gray, harris_keypoints1)
589 harris_descriptors1, harris_descriptors2 = compute_orb_descriptors(image2_gray, harris_keypoints2)
590
591
592 fast_descriptors1, fast_descriptors2 = compute_orb_descriptors(image1_gray, fast_keypoints1)
593 fast_descriptors2, fast_descriptors1 = compute_orb_descriptors(image2_gray, fast_keypoints2)
594
595
596
597
598 visualize_keypoints(image1, harris_keypoints1, "Harris Keypoints - School")
599 visualize_keypoints(image2, harris_keypoints2, "Harris Keypoints - Pixelated")
600
601
602 visualize_keypoints(image1, fast_keypoints1, "FAST Keypoints - School")
603 visualize_keypoints(image2, fast_keypoints2, "FAST Keypoints - Pixelated")
604
605
606
607
608
609 # Feature Matching (SSD & Ratio Test) for both Harris & FAST
610 matches_ssd_harris_1 = match_features_ssd(harris_descriptors1, harris_descriptors2)
611 matches_ratio_harris_1 = match_features_ratio(harris_descriptors1, harris_descriptors2)
612
613 matches_ssd_fast_1 = match_features_ssd(fast_descriptors1, fast_descriptors2)
614 matches_ratio_fast_1 = match_features_ratio(fast_descriptors1, fast_descriptors2)
615
616
617 # Visualize Results
618 visualize_matches(image1, harris_keypoints1, image2, harris_keypoints2, matches_ssd_harris_1, "
619     Harris + ORB (SSD) with Bernie darker")
620 visualize_matches(image1, harris_keypoints1, image2, harris_keypoints2, matches_ratio_harris_1, "
621     Harris + ORB (Ratio Test) with Bernie darker")
622
623 visualize_matches(image1, fast_keypoints1, image2, fast_keypoints2, matches_ssd_fast_1, "FAST + "
624     ORB (SSD) with Bernie darker")

```

```
622 | visualize_matches(image1, fast_keypoints1, image2, fast_keypoints2, matches_ratio_fast_1, "FAST+  
    ORB(Ratio Test)with Bernie darker")
```